CODE WALKTHROUGH

From a raw ASTM/HL7 byte stream to a row in lab_results

Every excerpt below is read out of the repository when this PDF is built, so the document always matches the code that ships. Bengali project overview: docs/PROJECT_GUIDE_BN.pdf.

MAIN.PY

1 · Entry point and wiring

`main.py` loads `config.yaml`, decides between Oracle and the in-memory store, then starts one listener per enabled analyzer. Serial analyzers get a thread, TCP analyzers get an asyncio task.

```
def start_listeners(config: Dict[str, Any], store: Any) -> List[Any]:
    started: List[Any] = []
    loop = asyncio.get_running_loop()

    for analyzer in config.get("analyzers", []) or []:
        if not analyzer.get("enabled", True):
            logger.info("skipped (disabled): %s", analyzer.get("name"))
            continue

        engine = build_engine(config, store, analyzer)
        conn_type = str(analyzer.get("connection_type", "TCP")).upper()

        if conn_type == "SERIAL":
            listener = SerialListener(analyzer, engine, loop=loop)
            listener.start()
        else:
            listener = TCPListener(analyzer, engine)
            asyncio.create_task(listener.start(), name=analyzer.get("name"))

        started.append(listener)

    return started


# ---------------------------------------------------------------------- main
async def run(config_path: str, dry_run: bool) -> None:
```

CORE/PARSER_FACTORY.PY

2 · Protocol sniffing

Analyzers get misconfigured more often than people admit. The sniffer looks at the first segment of the payload; `parse_message` retries with the other parser if the configured one returns nothing, and logs that the config is wrong.

```
def sniff_protocol(raw: str) -> str:
    """Return 'HL7' or 'ASTM' for a raw message payload."""
    if not raw:
        return "ASTM"
    text = raw.replace("\x0b", "").replace("\x1c", "").lstrip()
    first = text.split("\r")[0].split("\n")[0].strip()
    if any(first.upper().startswith(seg + "|") for seg in HL7_SEGMENTS):
        return "HL7"
    if first[:1].upper() == "H":
        return "ASTM"
    if "|" in text and any(f"|{seg}|" in text for seg in HL7_SEGMENTS):
        return "HL7"
    return "ASTM"


def build_parser(protocol: str, test_code_map: Optional[Dict[str, str]] = None):
    if protocol.upper() == "HL7":
        return HL7Parser(test_code_map)
```

```
def parse_message(
    raw: str,
    protocol: str = "AUTO",
    analyzer_name: str = "UNKNOWN",
    test_code_map: Optional[Dict[str, str]] = None,
) -> List[Dict[str, Any]]:
    """Parse with the configured protocol, falling back to sniffing."""
    chosen = protocol.upper()
    if chosen in ("AUTO", ""):
        chosen = sniff_protocol(raw)

    parser = build_parser(chosen, test_code_map)
    results = parser.parse(raw, analyzer_name)

    if not results and chosen != "AUTO":
        other = "HL7" if chosen == "ASTM" else "ASTM"
        retry = build_parser(other, test_code_map).parse(raw, analyzer_name)
        if retry:
            logger.warning(
                "%s sent %s but parsed as %s — check analyzer config",
                analyzer_name, chosen, other,
            )
            return retry

    return results
```

PARSERS/ASTM_PARSER.PY

3 · ASTM record parsing

Records are split on CR (LF and CRLF tolerated), then on |. Patient and order records are read once and attached to every result record, so each row is self-contained for the database.

```
def _split_records(raw: str) -> List[List[str]]:
    """Split a message into records, each record into fields."""
    records: List[List[str]] = []
    for line in _clean(raw).split(CR):
        line = line.strip()
        if not line:
            continue
        records.append([field.strip() for field in line.split("|")])
    return records


def _field(fields: List[str], index: int, default: str = "") -> str:
```

```
    def parse(self, raw_message: str, analyzer_name: str = "UNKNOWN") -> List[Dict[str, Any]]:
        records = _split_records(raw_message)
        if not records:
            logger.warning("empty ASTM message from %s", analyzer_name)
            return []

        header = self._header(records)
        sample = self._order(records)
        patient = self._patient(records)

        results: List[Dict[str, Any]] = []
        for fields in records:
            if not fields or fields[0].upper() != "R":
                continue
            result = self._result_record(fields, analyzer_name, header, sample, patient, raw_mes
    sage)
            if result:
                results.append(result)

        logger.info("%s: parsed %d result(s) from %s", self.name, len(results), analyzer_name)
        return results

    # ------------------------------------------------------------------ private
    def _header(self, records: List[List[str]]) -> Dict[str, str]:
        for fields in records:
```

PARSERS/HL7_PARSER.PY

4 · HL7 OBX parsing

Segments are grouped by name and the parser anchors on `OBX`. Field indexes follow the HL7 numbering (OBX-2 value type, OBX-3 identifier, OBX-5 value, OBX-8 flag).

```
    def _obx(
        self,
        fields: List[str],
        analyzer_name: str,
        context: Dict[str, Any],
        raw_message: str,
    ) -> Optional[Dict[str, Any]]:
        identifier = self._at(fields, 3)          # CODE^TEXT^LOINC
        value_type = self._at(fields, 2) or "NM"
        raw_value = self._at(fields, 5)

        code = identifier.split("^")[0].strip()
        if not code:
            return None
        code = self.test_code_map.get(code.upper(), code)

        abnormal = (self._at(fields, 8) or "N").upper()
        status = self._at(fields, 11) or "F"
        observed_at = _parse_hl7_datetime(self._at(fields, 14))

        return {
            "analyzer_name": analyzer_name,
            "sample_id": context.get("sample_id") or context.get("patient_id", ""),
            "patient_id": context.get("patient_id", ""),
            "patient_name": context.get("patient_name", ""),
            "test_code": code,
            "result_value": _to_float(raw_value) if value_type.upper() != "ST" else None,
            "result_text": raw_value,
            "result_unit": self._at(fields, 6),
            "ref_range": self._at(fields, 7),
            "abnormal_flag": abnormal,
            "result_status": status,
            "result_datetime": observed_at or context.get("sent_at") or datetime.now().strftime(
    "%Y-%m-%d %H:%M:%S"),
            "raw_message": raw_message,
```

CORE/VALIDATOR.PY

5 · Validation rules

One place decides what a valid result is: required keys, unit/flag normalisation, critical detection (`C`, `HH`, `LL`) and the control-sample filter.

```
    def normalise(self, result: Dict[str, Any]) -> Dict[str, Any]:
        out = dict(result)
        out["sample_id"] = str(out.get("sample_id") or "").strip()
        out["patient_id"] = str(out.get("patient_id") or "").strip()
        out["patient_name"] = str(out.get("patient_name") or "").strip()
        out["test_code"] = str(out.get("test_code") or "").strip().upper()
        out["analyzer_name"] = str(out.get("analyzer_name") or "").strip()
        out["result_unit"] = str(out.get("result_unit") or "").strip()
        out["ref_range"] = str(out.get("ref_range") or "").strip()
        out["result_text"] = str(out.get("result_text") or "").strip()

        flag = str(out.get("abnormal_flag") or "").strip().upper()
        out["abnormal_flag"] = FLAG_ALIASES.get(flag, flag or "N")

        out["is_critical"] = "Y" if out["abnormal_flag"] in CRITICAL_FLAGS else "N"
        out["is_abnormal"] = "Y" if out["abnormal_flag"] not in {"N", ""} else "N"

        if not out.get("result_datetime"):
            out["result_datetime"] = datetime.now().strftime("%Y-%m-%d %H:%M:%S")

        # numeric results may still arrive as strings from a custom parser
        if out.get("result_value") is not None:
            try:
                out["result_value"] = float(str(out["result_value"]).replace(",", "."))
            except (TypeError, ValueError):
                out["result_value"] = None

        out["protocol"] = str(out.get("protocol") or "").upper()
        return out
```

```
    def _problem(self, result: Dict[str, Any]) -> Optional[str]:
        for key in REQUIRED_FIELDS:
            if not result.get(key):
                return f"missing {key}"
        if result.get("result_value") is None and not result.get("result_text"):
            return "no value"
        if self.ignore_prefixes and result["sample_id"].upper().startswith(self.ignore_prefixes)
    :
            return f"control sample ({result['sample_id']})"
        if self.known_test_codes and result["test_code"] not in self.known_test_codes:
            return f"unknown test code {result['test_code']}"
        return None
```

CORE/ENGINE.PY

6 · The pipeline itself

Order matters: raw payload first, parse second, insert third, and only then build the ACK. If any step fails, the raw message is already safe and the analyzer gets a NAK instead of a silent drop.

```
    def process(self, raw: str, analyzer: Dict[str, Any], peer: str = "n/a") -> Dict[str, Any]:
        """Synchronous core — easy to unit test, easy to call from a thread."""
        analyzer_name = analyzer.get("name", "UNKNOWN")
        configured = str(analyzer.get("protocol", "AUTO")).upper()
        started = datetime.now()

        message_id = self.store.log_raw_message(analyzer_name, configured, raw, peer)
        self.counters["messages"] += 1
```

```
        entry: Dict[str, Any] = {
            "message_id": message_id,
            "analyzer_name": analyzer_name,
            "peer": peer,
            "received_at": started.strftime("%Y-%m-%d %H:%M:%S"),
            "configured_protocol": configured,
            "raw": raw,
        }

        try:
            protocol = sniff_protocol(raw)
            entry["detected_protocol"] = protocol

            parsed = parse_message(
                raw,
                protocol=configured if configured != "AUTO" else "AUTO",
                analyzer_name=analyzer_name,
                test_code_map=self.test_code_map,
            )
            entry["parsed_count"] = len(parsed)

            if not parsed:
                entry.update(status="ERROR", note="no results parsed", stored=0)
                self.counters["errors"] += 1
                self.store.mark_message(message_id, "ERROR", "no results parsed")
                self._record(entry)
                return self._response(False, entry, "no results parsed")

            validated = self.validator.validate(parsed)
            entry["valid_count"] = len(validated)
            if not validated and getattr(self.validator, "last_ignored", None):
                sample = getattr(self.validator, "last_ignored")[0].split("(")[-1].rstrip(")")
                note = f"control sample {sample} ignored (not patient data)"
                entry.update(status="IGNORED", note=note, stored=0)
                self.store.mark_message(message_id, "IGNORED", note)
                self.store.insert_audit("IGNORED", analyzer_name, note)
                self._record(entry)
                return self._response(True, entry, note)

            if not validated:
                entry.update(status="ERROR", note="all results failed validation", stored=0)
                self.counters["errors"] += 1
                self.store.mark_message(message_id, "ERROR", "validation failed")
                self._record(entry)
                return self._response(False, entry, "validation failed")

            stored = self.store.insert_results(validated, raw)
            duplicates = len(validated) - stored
            self.counters["results"] += stored
            self.counters["duplicates"] += duplicates

            status = "DUPLICATE" if stored == 0 else "PARSED"
            note = f"{stored} stored" + (f", {duplicates} duplicate(s)" if duplicates else "")
```

DATABASE/DB_WRITER.PY

7 · Oracle writer

Two tables: `lis_raw_messages` (verbatim archive, with status) and `lab_results` (one row per test, de-duplicated). The duplicate check is cheap and keeps the analyzer's re-transmissions out of the report.

```
    def is_duplicate(self, result: Dict[str, Any]) -> bool:
        sql = f"""
            SELECT COUNT(*)
              FROM {self.schema}.lab_results
             WHERE sample_id     = :sample_id
               AND test_code     = :test_code
               AND analyzer_name = :analyzer_name
               AND result_value  = :result_value
        """
        with self.db.cursor() as cur:
            cur.execute(
                sql,
                {
                    "sample_id": result["sample_id"],
                    "test_code": result["test_code"],
                    "analyzer_name": result["analyzer_name"],
                    "result_value": result.get("result_value"),
                },
            )
            (count,) = cur.fetchone()
```

```
    def insert_results(self, results: Iterable[Dict[str, Any]], raw_message: str = "") -> int:
        """Insert results, skipping duplicates. Returns how many were stored."""
        stored = 0
        for result in results:
            try:
                if self.is_duplicate(result):
                    logger.warning(
                        "duplicate skipped: %s / %s",
                        result["sample_id"], result["test_code"],
                    )
                    continue
            except Exception as exc:
                logger.error("duplicate check failed for %s: %s", result.get("test_code"), exc)
                continue
```

LISTENERS/TCP_LISTENER.PY

8 · Listener framing

`_extract` turns a byte stream into messages: MLLP (`\x0b…\x1c`), ASTM (`\x05`/`\x03`), or bare text closed by the client. Whatever framing arrives, the analyzer gets `0x06` for success and `0x15` for a rejected message.

```
    def _extract(self, buffer: bytes) -> tuple[Optional[bytes], bytes, bool]:
        """Pull one complete message out of the stream buffer.

        Returns ``(payload, remaining, needs_ack)``. ``payload`` is None when
        no complete message is available yet.
        """
        if not buffer:
            return None, buffer, False

        # drop stray bytes analyzers leave between frames (EOT / NUL / ENQ)
        buffer = buffer.lstrip(b"\x04\x00\x05")
        if not buffer:
            return None, buffer, False

        # --- HL7 MLLP: \x0b ... \x1c ------------------------------
        start = buffer.find(MLLP_START)
        if start != -1:
            end = buffer.find(MLLP_END, start)
            if end == -1:
                return None, buffer, False
            payload = buffer[start + 1:end]
            remaining = buffer[end + 1:]
            if remaining.startswith(CR):
                remaining = remaining[1:]
            return payload, remaining, True
        if buffer.startswith(b"\x0b"):
            return None, buffer, False

        # --- ASTM: \x05 "H|..." \x02 ... \x03 ----------------------
        if buffer.startswith(b"\x05"):  # ENQ — strip it, message follows
            buffer = buffer[1:]

        if MLLP_END in buffer:  # stray FS from a mis-framed client
            idx = buffer.find(MLLP_END)
            payload, remaining = buffer[:idx], buffer[idx + 1:]
            return payload, remaining, True

        if b"\x03" in buffer:  # ASTM ETX terminates the record set
            idx = buffer.find(b"\x03")
            payload, remaining = buffer[:idx], buffer[idx + 1:]
            return payload, remaining, True
```

TESTS/TEST_PIPELINE.PY

9 · Proving it works

Twenty-eight tests cover the parsers, the sniffer, the validator, de-duplication, the ACK modes and the error paths. No Oracle, no sockets, no lab instrument needed.

test run

```
python -m unittest discover -s tests -v
  ...
  Ran 28 tests in 0.006s
  OK
```

```
    def test_end_to_end_astm(self):
        response = self.engine.process(SAMPLES["astm_cbc_normal"], ANALYZER, peer="tcp:1")
        self.assertTrue(response["ok"])
        self.assertEqual(response["status"], "PARSED")
        self.assertEqual(response["stored"], 4)
        self.assertEqual(response["ack"], "\x06")
        # HGB was remapped to HB by the config map
        stored_codes = {r["test_code"] for r in self.store.results}
        self.assertIn("HB", stored_codes)
        self.assertEqual(self.store.stats()["total"], 4)

    def test_end_to_end_hl7(self):
```

Where to start reading

Follow the order of this document: main.py → engine.py → parsers/ → validator.py→ db_writer.py. Everything else (demo, tools) reuses those five files.